

Fusion Model Release Control Procedure

Purpose

This procedure defines how the Fusion Model cell handles candidate observations during the Northern Veil planning exercise. The release threshold is **0.7342** before a candidate observation can enter the shared planning workspace. The cell exists to turn exercise reports, simulated sensor notes, staff updates, and controller injects into a controlled planning signal.

The procedure does not rely on a single source stream. It assumes that reports arrive with uneven quality, inconsistent timing, and different access labels. The Fusion Model helps staff compare those inputs without treating the score as a final truth claim. The score is a workflow control. It supports release, hold, withdrawal, and review decisions.

The controlling release value for this procedure is the confidence threshold **0.7342**. A candidate observation must meet or exceed **0.7342** before it can be released into the shared planning workspace, unless a watch lead records a named exception under the exception rules below.

Operating Frame

The Fusion Model receives four input classes during the exercise. Simulated collection reports describe field observations generated by the scenario. Staff situation updates describe planning assumptions that have already reached a coordination desk. Logistics availability notes describe notional constraints on fuel, maintenance, lift, and staging capacity. Controller injects introduce scenario changes that test whether the staff can adjust without widening access beyond the intended audience.

Each input receives three component values. Freshness estimates whether the item is recent enough to affect a current planning decision. Consistency estimates whether the item aligns with other reports that describe the same exercise event. Relevance estimates whether the item changes a decision rather than merely adding background detail. The model combines those values into a confidence score used by the watch team.

The Fusion Model does not override access labels. A high-confidence item with a mismatched access label remains blocked. A low-confidence item may still matter, but it cannot become a released observation without a documented exception. Analysts must treat the model as a structured review aid, not as an authority to distribute information.

Observation States

Draft observations remain inside analyst review. A draft may contain unresolved source conflicts, missing timing, or unclear relevance. Draft material should not be copied into planning products. Analysts may discuss a draft inside the cell to improve the record, but they must not use it as evidence for external coordination.

Candidate observations have enough structure to be reviewed by the watch lead. A candidate normally includes at least two input classes, a named planning decision, an access label, and a plain-language caveat.

The candidate still requires a confidence score at or above **0.7342** before normal release. Candidate status means the observation is ready for release review. It does not mean planners can act on it.

Released observations may enter the shared planning workspace. A released observation must include the confidence value, the release rationale, the access label, the caveat, and the receiving workspace. The watch lead must record whether the release followed the normal threshold or used an exception.

Withdrawn observations are removed from planning use. Withdrawal may happen because a later report supersedes the item, because the access label was wrong, because the relevance was overstated, or because the scenario controller changed the assumption. Withdrawn observations remain in the exercise record for review.

Release Review

The release review begins when an analyst promotes a candidate observation. The watch lead checks the input count, the named planning decision, the access label, the confidence score, and the caveat. The lead should read the underlying notes before approving release. The score helps prioritize review, but the lead remains responsible for the final release action.

A candidate observation that meets or exceeds **0.7342** can move to released status when the access label matches the receiving workspace and no source conflict remains unresolved. If the score is below **0.7342**, the watch lead can still approve an exception only when the simulated safety or continuity value outweighs the delay. The exception must include the below-threshold score and a follow-up review time.

The lead must block a candidate observation even above **0.7342** when the access label is unclear, the receiving workspace is not allowed to see the item, the caveat is missing, or the underlying inputs describe different events. Blocking is not an analytic rejection. It is a routing protection while staff correct the record.

Caveats

Every released observation includes a caveat. The caveat states whether the item is time-sensitive, whether it depends on one high-weight input, whether it supersedes an earlier item, and whether staff should expect a follow-up. The caveat should help planners act with the right level of confidence. It should not reproduce the full analytic chain.

Caveats use direct language. A good caveat says, "Use for staging decisions until the next watch review," or "Do not use for final allocation because logistics availability is unresolved." A poor caveat says, "Model confidence is acceptable," because that does not tell the planner what the limits are.

If a released observation supersedes an earlier observation, the earlier entry is marked as superseded. Staff must not delete the earlier entry. Planners need the change history to understand why a decision moved. Reviewers need the chain to understand whether the model was too slow, too fast, or too sensitive to one input.

Exception Rules

The watch lead may release an observation below **0.7342** when a simulated safety condition requires immediate planning attention. The exception record must list the score, the reason, the receiving audience, the caveat, and the time for follow-up. The exception applies only to that observation. It does not lower the threshold for future releases.

The watch lead may also suspend a release above **0.7342** when the access label does not match the destination, when the source conflict affects the planning decision, or when a controller inject creates unresolved authority. The suspension record must name the issue and the condition for restart. The analyst who created the candidate should correct the issue before the next watch review.

Exception patterns are reviewed at the end of each exercise day. Repeated below-threshold release may mean the model is too conservative, the exercise injects are poorly timed, or staff are using the exception path as a shortcut. Repeated high-score blocking may mean metadata quality is weak or the receiving workspace has unclear access boundaries.

Review Products

The cell maintains a release register. Each row includes the observation title, input classes used, confidence score, access label, release state, caveat, receiving workspace, lead decision, and review time. The register is the main record for after-action review. It allows staff to compare model output with actual release behavior.

The daily review summarizes released items, held items, withdrawals, exceptions, and blocked high-score candidates. The review should separate scoring issues from workflow issues. A scoring issue means the model did not represent the exercise condition well. A workflow issue means staff applied the access label, caveat, or release state incorrectly.

The after-action review should not reward high release volume. The goal is useful, authorized planning evidence. A lower volume of clear released observations is better than a large set of weak signals that planners cannot safely use.